

SAiDL Summer 2026 Induction Assignment

Long-Context Sequence Modelling

An Empirical Study of Attention Variants, Positional Encodings,
Convolutional Hybrids, and Attention-Free Transformers on WikiText-2

Core ML Track

Aditi Kharya

June 2026

Contents

1	Introduction	4
1.1	Foundations: Parameters, Gradients, and Perplexity	4
1.2	The Transformer Architecture	4
1.3	Scaled Dot-Product Attention	4
1.4	The Quadratic Bottleneck and Research Motivation	5
2	Experimental Setup	5
2.1	Dataset and Tokenisation	5
2.2	Compute Environment and State Persistence Constraints	5
2.3	Base Model Configuration	6
2.4	Hardware and Software	6
2.5	Evaluation Metrics and Metric Scope	6
2.6	Requirements Traceability	7
3	Task 1: Baseline Transformer	7
3.1	Architecture	7
3.2	Training Stability	7
3.3	Results	7
4	Task 2: Attention Variants	8
4.1	Motivation	8
4.2	Sliding Window Attention	8
4.3	Multi-Query Attention (MQA)	8
4.4	Linear Attention	9
4.5	Results at Context Length 1024	9
4.6	Multi-Context Length Evaluation	10
4.7	Analysis	11
5	Task 3: Positional Encodings	12

5.1	Limitations of Learned Absolute Embeddings	12
5.2	Rotary Positional Embedding (RoPE)	12
5.3	ALiBi (Attention with Linear Biases)	12
5.4	Relative Positional Encoding	13
5.5	Results at Context 1024	13
5.6	Extrapolation Test	14
5.7	Analysis	14
6	Task 4: Convolution + Attention Hybrids	15
6.1	Motivation	15
6.2	Design 1: Causal Conv1D Before Each Attention Block	15
6.3	Design 2: Interleaved Conv and Attention Blocks	16
6.4	Best-Variant Combined Models	16
6.5	Results	17
6.6	Analysis	17
7	Bonus: Attention-Free Transformer (AFT) Variants	18
7.1	Mechanism	18
7.2	Implementation Note	18
7.3	Experimental Constraints for AFT Evaluation	19
7.4	Results	19
7.5	Analysis	19
8	Task 5: Summary and Comparative Analysis	20
8.1	Hardware Profiling	20
8.2	Master Comparison Table	21
8.3	Key Findings	21
9	Limitations	22
10	Conclusion	23

Abstract

This report documents a systematic empirical study of Transformer architectural modifications for long-context language modelling, trained on WikiText-2. Starting from a standard decoder-only Transformer baseline (validation perplexity 406.68 at context length 1024), we evaluate three families of modifications: attention variants (Sliding Window, Multi-Query Attention [2], Linear Attention [3]), positional encodings (RoPE [4], ALiBi [5], Relative PE [6]), and convolutional hybrids. We additionally implement the four Attention-Free Transformer (AFT) variants from Zhai et al. [10] as a bonus. The central finding is that *positional encoding choice dominates model quality* at this scale and training budget: RoPE and ALiBi reduce perplexity by over 40% relative to the baseline, far outpacing any attention variant. Combining causal Conv1D, Sliding Window attention, and RoPE yields the strongest model at perplexity 225.55. Extrapolation tests confirm ALiBi as the most robust encoding for length generalisation, with perplexity improving monotonically at sequences longer than its training context.

All experiments were conducted on free-tier Kaggle T4 GPUs (15 GB VRAM) under a fixed 2500-step training budget. The report is fully transparent about its constraints: throughput metrics reflect *training* processing speed (with gradients), not inference latency; batch sizes were reduced at context 2048 due to memory limits; all runs use a single seed; and model checkpoints stored in the ephemeral `/kaggle/working/` directory were lost to session timeouts — all reported metrics were extracted and persisted locally before environment teardown. Relative comparisons are the primary takeaway; absolute perplexity values should not be treated as benchmark-quality numbers.

1 Introduction

1.1 Foundations: Parameters, Gradients, and Perplexity

Deep learning learns parametric mappings from data by composing affine transformations with pointwise non-linearities. For layer l :

$$\mathbf{z}^{(l)} = W^{(l)}\mathbf{x}^{(l-1)} + \mathbf{b}^{(l)}, \quad \mathbf{x}^{(l)} = \sigma(\mathbf{z}^{(l)}) \quad (1)$$

Without σ , any stack of linear layers collapses to a single affine map. Modern language models use the Gaussian Error Linear Unit (GELU) [9]:

$$\text{GELU}(x) = x \cdot \Phi(x), \quad \Phi(x) = P(X \leq x), \quad X \sim \mathcal{N}(0, 1) \quad (2)$$

GELU is preferred over ReLU because small negative activations retain a gradient signal (avoiding dying-neuron pathologies), the function is infinitely differentiable everywhere, and the probabilistic gating acts as a soft regulariser during training.

For autoregressive language modelling the objective is cross-entropy over token sequences:

$$\mathcal{L}(\Theta) = -\frac{1}{N} \sum_{i=1}^N \sum_{k=1}^V y_{i,k} \log \hat{y}_{i,k} \quad (3)$$

Perplexity = $\exp(\mathcal{L}_{\text{val}})$. A perplexity of P means the model is as uncertain at each token as if it were choosing uniformly among P equally likely options. All comparisons in this report use validation perplexity as the primary quality metric. Parameters are updated with AdamW and gradient clipping (norm 1.0) at every step.

1.2 The Transformer Architecture

The Transformer [1] processes sequences in parallel via self-attention. Tokens are embedded and summed with positional information:

$$H^{(0)} = X_{\text{tok}} + X_{\text{pos}} \quad (4)$$

Each of L blocks applies *Pre-LayerNorm* residual connections (LN before each sub-block rather than after, which produces more stable gradients at small model sizes):

$$\tilde{H}^{(l)} = H^{(l-1)} + \text{MHA}(\text{LN}(H^{(l-1)})) \quad (5)$$

$$H^{(l)} = \tilde{H}^{(l)} + \text{FFN}(\text{LN}(\tilde{H}^{(l)})) \quad (6)$$

1.3 Scaled Dot-Product Attention

$$\text{Attention}(Q, K, V) = \text{softmax}\left(\frac{QK^\top}{\sqrt{d_k}} + M\right)V \quad (7)$$

where causal mask $M_{ij} = -\infty$ for $j > i$ and 0 otherwise. Multi-head attention runs h parallel heads at dimension $d_k = d_v = d_{\text{model}}/h$.

1.4 The Quadratic Bottleneck and Research Motivation

Computing $QK^\top$ requires $\mathcal{O}(T^2 d_{\text{model}})$ work and $\mathcal{O}(hT^2)$ memory. This quadratic cost is directly observable in our experiments: at context length 2048, standard full attention required a reduced batch size (see Section 9), while efficient variants handled larger batches. The KV-cache memory at inference scales as:

$$\text{KV bytes} = 2 \times L_{\text{layers}} \times h \times d_k \times T \times \text{bytes-per-element} \quad (8)$$

Beyond raw compute, learned absolute position tables are hard-bounded: a position index beyond T_{train} has no representation, causing out-of-distribution failures at inference time. This report addresses both challenges empirically.

2 Experimental Setup

2.1 Dataset and Tokenisation

All experiments use WikiText-2 [8], tokenised with the GPT-2 BPE tokeniser (vocabulary size 50,257), loaded via the HuggingFace `datasets` library with the `Salesforce/wikitext` path. The training set contains approximately **2,088,628 tokens** and the validation set approximately 217,646 tokens as measured by our tokenisation pipeline. Standard deviations σ reflect the empirical macro-variance observed in separate stochastic evaluation batch subsets during the validation phase.

2.2 Compute Environment and State Persistence Constraints

All model training sequences were executed on interactive Kaggle cloud instances provisioned with ephemeral NVIDIA T4 hardware accelerators under a fixed computational budget. Due to runtime session timeouts and automated container disconnections inherent to the free cloud tier, local checkpoint files (`.pt` files stored in the volatile `/kaggle/working/` directory) were lost upon environment teardown.

To maintain strict scientific transparency and avoid data fabrication, all evaluation metrics reported in our summary matrices represent the definitive validation values extracted and persisted into local summary objects *before* container recycling. Due to these platform-enforced constraints, all training runs are treated as single-seed executions, and the standard deviations (σ) reported throughout this document reflect the empirical variance observed *across stochastic evaluation batches* drawn during each validation call, not cross-seed variance. This quantifies how sensitive each architecture is to the particular token distribution of a given evaluation batch.

Context length 4096 was evaluated but triggered an Out-Of-Memory (OOM) allocation failure during the initial cache build on the T4 (15 GB VRAM), confirming 2048 tokens as our practical execution ceiling.

2.3 Base Model Configuration

Table 1: Shared hyperparameters across all experiments. Any deviation is noted explicitly in the relevant section.

Hyperparameter	Value
Number of layers	4
Embedding dimension	256
Number of attention heads	4
Head dimension	64
FFN expansion factor	$4\times$ (hidden dim = 1024)
Default context length	1024
Dropout	0.1
Default batch size	8
Learning rate	3×10^{-4}
Optimiser	AdamW
Gradient clipping norm	1.0
Training steps	2500
Evaluation interval	every 500 steps
Total parameters (baseline)	$\approx 29\text{M}$

2.4 Hardware and Software

All experiments ran on Kaggle Notebooks (free tier), single NVIDIA Tesla T4 GPU (15.36 GB VRAM, Compute Capability 7.5), 2-core Intel Xeon @ 2.00 GHz, 31.35 GB system RAM, Linux 6.6.122+, Python 3.12.13, PyTorch with CUDA 12.8, NVIDIA Driver 580.105.08. Two T4 GPUs were available in the Kaggle environment but all training ran on GPU 0 in isolation for consistent memory reporting.

2.5 Evaluation Metrics and Metric Scope

- **Validation perplexity** ($\mu \pm \sigma$) — primary quality metric, with σ across evaluation batches.
- **Peak GPU memory (MB)** — via `torch.cuda.max_memory_allocated()`, reset before each profiling window.
- **Training time per epoch (min)** — $(\text{total training time} / \text{max_steps}) \times (\text{train tokens} / \text{tokens per step})$.
- **Throughput (tokens/sec)** — cumulative tokens processed divided by elapsed wall-clock time, measured *during training with gradients active*. This is **training throughput**, not inference latency. True inference latency (autoregressive decoding without gradient computation) would be measured separately in a `torch.no_grad()` loop; we do not report inference latency here.

2.6 Requirements Traceability

Table 2: Task completion against the SAiDL Summer 2026 assignment specification.

Task	Status	Notes
1. Baseline Transformer	Complete	Context 1024, perplexity 406.68
2. Attention Variants (≥ 3)	Complete	SW, MQA, Linear; multi-context eval
3. Positional Encodings	Complete	RoPE, ALiBi, Relative PE; extrapolation test
4. Conv Hybrids	Complete	2 basic designs + 4 combined best-variant models
5. Logging & Reporting	Complete	GPU memory, epoch time, throughput, master table
Bonus: AFT Variants	Complete	AFT-Simple/Full/Local/Conv at context 256

3 Task 1: Baseline Transformer

3.1 Architecture

We implement a standard decoder-only Transformer following the architecture of Vaswani et al. [1]. Each block uses Pre-LN residual connections, causal masked multi-head self-attention, and a two-layer FFN with GELU activation at $4\times$ hidden expansion. Position information is encoded via a learned absolute embedding lookup table $W_{\text{pos}} \in \mathbb{R}^{T_{\text{max}} \times d_{\text{model}}}$.

The causal mask is registered as a persistent buffer to avoid GPU reallocation at every forward pass:

```
self.register_buffer("mask",  
    torch.tril(torch.ones(context_len, context_len))  
        .view(1, 1, context_len, context_len))
```

Future positions are masked to $-\infty$ before softmax, enforcing the autoregressive constraint.

3.2 Training Stability

Training loss decreased smoothly and monotonically across all 2500 steps with no erratic gradient spikes observed. Gradient clipping at norm 1.0 was active throughout but rarely saturated after step 100. The throughput jump from ≈ 285 tokens/sec at step 0 to $\approx 17,000$ tokens/sec by step 500 reflects GPU JIT compilation and CUDA kernel caching warming up, not training instability. No NaN losses or divergences occurred in any run.

3.3 Results

Table 3: Baseline Transformer performance at context length 1024, batch size 8.

Metric	Value
Validation perplexity ($\mu \pm \sigma$)	406.68 ± 17.42
Peak GPU memory	8,346.2 MB (8.15 GB)
Total training time (2500 steps)	19.33 min
Avg. training throughput	17,657 tokens/sec
Parameters	29,203,537

Key finding. The baseline delivers perplexity 406.68 at context 1024. Because all subsequent models share identical hyperparameters and step count, differences in perplexity are attributable to the architectural modification, not training variance or scheduling differences.

4 Task 2: Attention Variants

4.1 Motivation

Standard self-attention builds an explicit $T \times T$ score matrix whose $\mathcal{O}(T^2d)$ cost and $\mathcal{O}(hT^2)$ memory becomes prohibitive for long sequences. We evaluate three alternatives that trade exactness for different forms of efficiency, and examine how each scales as context length grows from 512 to 2048.

4.2 Sliding Window Attention

Mechanism. Each token i attends only to a fixed window w of preceding positions:

$$M_{ij} = \begin{cases} 0 & \text{if } \max(0, i-w) \leq j \leq i \\ -\infty & \text{otherwise} \end{cases} \quad (9)$$

We use $w = 64$. Complexity falls from $\mathcal{O}(T^2)$ to $\mathcal{O}(T \cdot w) = \mathcal{O}(T)$ since w is fixed. The hypothesis: at this small model scale, local syntactic dependencies dominate, and global attention wastes capacity on irrelevant distant tokens.

Implementation. The mask is computed with vectorised distance comparisons — no Python for-loop:

```
positions = torch.arange(T, device=x.device)
dist = positions.unsqueeze(0) - positions.unsqueeze(1)
sw_mask = torch.where(
    (dist >= -self.window_size) & (dist <= 0),
    torch.zeros(T, T, device=x.device),
    torch.full((T, T), float('-inf'), device=x.device))
```

Why Sliding Window uses more memory than expected. Despite the smaller attention window, peak memory at context 1024 (10,566 MB) exceeds the baseline (8,346 MB). The reason is purely implementational: we materialise the full $T \times T$ attention score tensor and then apply the mask, so PyTorch allocates the full matrix before any values are zeroed. A sparse kernel (e.g. FlashAttention’s local-window path) would avoid this entirely. This is an *implementation* limitation, not an architectural one. The theoretical memory footprint of Sliding Window Attention is genuinely smaller.

4.3 Multi-Query Attention (MQA)

Mechanism. Standard MHA maintains h independent K and V projections. MQA [2] shares a single K and V across all heads while keeping head-specific Q projections:

$$\text{head}_g = \text{Attention}\left(XW_Q^{(g)}, XW_K, XW_V\right) \quad (10)$$

KV parameter count shrinks from $\mathcal{O}(h \cdot d_{\text{model}} \cdot d_k)$ to $\mathcal{O}(d_{\text{model}} \cdot d_k)$. With 4 heads and $d_k = 64$, the KV activation footprint is $4\times$ smaller per layer, giving this model 28,808,785 parameters ($\approx 400\text{K}$ fewer than baseline).

Why MQA runs at larger batches at context 2048. The per-step KV activation tensors are $4\times$ smaller per layer. At context 2048 with batch 4, this saves approximately $4 \times (2048 \times 64 \times 4 \text{ layers}) \approx 2 \text{ GB}$ of intermediate activations, enough to fit comfortably within the T4 budget. This is a direct consequence of shared KV heads, not any training trick.

Why MQA underperforms baseline on perplexity. Sharing a single KV head means each of the 4 query heads must compete for information from the same key/value subspace. At $d_k = 64$, this is a real expressivity bottleneck: different heads cannot independently specialise to extract syntactic versus semantic patterns from the same context. MQA’s practical advantage is in *inference* (KV-cache reduction during autoregressive decoding), not in training quality, and this advantage grows at scale. Neither of which we test here.

4.4 Linear Attention

Mechanism. Linear Attention [3] replaces the softmax with a kernel function $\phi(x) = \text{ELU}(x) + 1$, enabling computation reordering via the associative property of matrix multiplication:

$$\text{output}_t = \frac{\phi(q_t)^\top \sum_{s \leq t} \phi(k_s) v_s^\top}{\phi(q_t)^\top \sum_{s \leq t} \phi(k_s) + \epsilon} \quad (11)$$

By computing $K^\top V$ first — a fixed $d \times d$ matrix regardless of T — the $T \times T$ score matrix is never materialised. Complexity drops to $\mathcal{O}(Td^2)$, i.e. $\mathcal{O}(T)$ since $d = 64$ is constant.

Why ELU+1? Softmax guarantees positive, normalised attention weights. Replacing it with $\phi(x) = \text{ELU}(x) + 1$ preserves positivity (ELU output ≥ -1 , so $+1$ gives output ≥ 0) without requiring the full row of scores to exist first. This is precisely what allows the bracket-reordering trick.

4.5 Results at Context Length 1024

Table 4: Attention variant comparison at context length 1024, batch size 8. σ is empirical variance across evaluation batches.

Model	PPL ($\mu \pm \sigma$)	Peak GPU (MB)	Training Throughput	Complexity
Baseline (MHA)	406.68 $\pm$ 17.42	8,346	17,657 tok/s	$\mathcal{O}(T^2)$
Sliding Window	385.78 $\pm$ 13.10	10,566	17,554 tok/s	$\mathcal{O}(Tw)^*$
MQA	417.42 $\pm$ 23.05	8,341	17,790 tok/s	$\mathcal{O}(T^2)$
Linear Attention	389.27 $\pm$ 31.45	7,367	20,602 tok/s	$\mathcal{O}(T)$

*Dense-mask implementation; theoretical memory is lower with a sparse kernel.





Figure 1: Attention variant perplexity at context 1024. Error bars show empirical σ across evaluation batches.


4.6 Multi-Context Length Evaluation

Table 5: Attention variants across context lengths. Batch sizes at context 2048: Baseline, Sliding Window, MQA, and Linear Attention all ran at batch size 4 (memory cleared between runs allowed this; see Section 9). Context length 4096 triggered OOM and was not evaluated.


Model	CTX	PPL ($\mu \pm \sigma$)	Batch	Peak GPU (MB)	Epoch (min)	Training Throughput
Baseline	512	434.80 $\pm$ 16.25	8	4,058	1.86	18,712 tok/s
Baseline	1024	406.68 $\pm$ 17.42	8	8,346	3.94	17,657 tok/s
Baseline	2048	448.05 $\pm$ 24.10	4	10,855	2.74	14,528 tok/s
Sliding Window	512	425.59 $\pm$ 12.04	8	6,292	2.23	17,870 tok/s
Sliding Window	1024	385.78 $\pm$ 13.10	8	10,566	2.27	17,554 tok/s
Sliding Window	2048	392.31 $\pm$ 18.55	4	11,561	2.92	13,646 tok/s
MQA	512	434.90 $\pm$ 21.30	8	4,053	2.10	19,028 tok/s
MQA	1024	417.42 $\pm$ 23.05	8	8,341	2.24	17,790 tok/s
MQA	2048	449.39 $\pm$ 29.80	4	11,637	2.71	14,714 tok/s
Linear Attn	512	402.36 $\pm$ 28.50	8	3,864	2.09	19,933 tok/s
Linear Attn	1024	389.27 $\pm$ 31.45	8	7,367	2.27	20,602 tok/s
Linear Attn	2048	435.99 $\pm$ 42.15	4	8,998	2.51	22,201 tok/s




(a) Validation perplexity vs. context length



(b) Peak GPU memory vs. context length



(c) Training throughput vs. context length



(d) Evaluation batch variance σ vs. context length

Figure 2: Multi-context comparison of four attention variants across perplexity, memory, throughput, and evaluation stability. Sliding Window is the most stable across lengths.

4.7 Analysis

Sliding Window is the best-quality attention variant (385.78 at 1024 vs. baseline 406.68), validating the local-context hypothesis. The perplexity curve $425.59 \rightarrow 385.78 \rightarrow 392.31$ across context lengths is the smoothest of the four models. By restricting each token to its 64 nearest predecessors, the model is forced to build maximally informative local representations rather than spreading limited capacity across all T tokens, many of which are semantically irrelevant to the immediate prediction. The slight rise at 2048 is attributable to the smaller batch size (4 vs. 8) introducing noisier gradient estimates, not an architectural failure.

Sliding Window’s counter-intuitive memory overhead is worth explaining carefully, because it appears to contradict the theoretical benefit. Our PyTorch implementation creates a dense $T \times T$ mask tensor, applies it to the full attention score matrix, and *then* softmaxes. Every element of $QK^\top$ is computed before any masking — the GPU allocates the full matrix first. A sparse implementation or FlashAttention-style blocked kernel would prevent this allocation entirely, realising the expected memory savings.

MQA underperforms baseline (417.42 vs. 406.68) because at $d_k = 64$ with only 4 heads,

forcing all heads to share a single KV projection is a measurable expressivity bottleneck. Heads that would specialise in syntactic vs. semantic patterns must compete for the same context vector. This penalty is well-documented to shrink at larger model scales (where d_k and depth are both larger), and MQA’s real benefit — dramatically smaller KV cache during autoregressive *inference* — is not captured by training throughput.

Linear Attention leads on memory and throughput at all context lengths. At context 2048, it uses 8,998 MB vs. the baseline’s 10,855 MB and delivers 22,201 tok/s vs. 14,528. This directly demonstrates the $\mathcal{O}(T)$ vs. $\mathcal{O}(T^2)$ memory advantage.

Batch-size confound at context 2048. After clearing GPU memory between runs, all four models including the Baseline ran at batch size 4 at context 2048. However, the gradient estimates at batch 4 are noisier than at batch 8, which partially accounts for the perplexity degradation at 2048 across all models. Memory comparisons at 2048 are the more reliable metric.

5 Task 3: Positional Encodings

5.1 Limitations of Learned Absolute Embeddings

The baseline uses a lookup table $W_{\text{pos}} \in \mathbb{R}^{T_{\text{max}} \times d}$ assigning a trainable vector to each absolute position index. Two fundamental problems:

1. **Hard boundary.** Position $T_{\text{max}} + 1$ has no representation. Hard failures occur at inference time for longer sequences.
2. **Absolute, not relative.** The model must infer from data that position 500 is “far from” position 1 — wasting parameter budget and training steps that could instead be spent on semantic content learning.

5.2 Rotary Positional Embedding (RoPE)

RoPE [4] rotates query and key vectors in 2D subspaces using position-dependent rotation matrices. For positions m and n :

$$\langle R_m q, R_n k \rangle = \langle R_{m-n} q, k \rangle \quad (12)$$

The inner product depends only on the *relative* offset $m - n$. For each dimension pair (x_{2i}, x_{2i+1}) in a head:

$$\begin{pmatrix} x'_{2i} \\ x'_{2i+1} \end{pmatrix} = \begin{pmatrix} \cos(t\theta_i) & -\sin(t\theta_i) \\ \sin(t\theta_i) & \cos(t\theta_i) \end{pmatrix} \begin{pmatrix} x_{2i} \\ x_{2i+1} \end{pmatrix}, \quad \theta_i = 10000^{-2i/d} \quad (13)$$

No additional trainable parameters. The key extrapolation property: the formula is defined for any t , so positions beyond T_{train} produce valid (if out-of-distribution) rotation angles rather than undefined lookups.

5.3 ALiBi (Attention with Linear Biases)

ALiBi [5] subtracts a linear distance penalty directly from attention logits:

$$\text{Attn}_{i,j} = \text{softmax} \left(\frac{q_i^\top k_j}{\sqrt{d_k}} - m_h \cdot |i - j| \right) \quad (14)$$

where $m_h = 2^{-8h/H}$ is a fixed head-specific slope. ALiBi has *zero* learned position parameters. The penalty is a deterministic function of the integer distance $|i - j|$. This is the most principled extrapolator: the rule is defined for any distance, and the linear penalty naturally downweights distant tokens as a built-in smoothing prior regardless of sequence length.

5.4 Relative Positional Encoding

Following Shaw et al. [6], we add learnable relative position biases to the attention scores:

$$A_{ij} = \frac{q_i^\top k_j}{\sqrt{d_k}} + \frac{q_i^\top r_{i-j}}{\sqrt{d_k}} \quad (15)$$

where $r_{i-j} \in \mathbb{R}^{d_k}$ is drawn from a table of size $2T_{\max} - 1$. This adds 29,465,425 parameters. The hard limitation: the table covers distances $[-(T - 1), +(T - 1)]$ for training length T . Evaluating on longer sequences produces out-of-range indices, causing a hard runtime failure.

For the extrapolation experiment, all three PE models were trained on $L_{\text{train}} = 512$ and evaluated at $L_{\text{test}} \in \{512, 1024, 2048\}$ without retraining.

5.5 Results at Context 1024

Table 6: Positional encoding comparison at context 1024, batch 8.

Model	PPL ($\mu \pm \sigma$)	Peak GPU (MB)	Epoch (min)	Training Throughput
Baseline (Learned Abs.)	406.68 ± 17.42	8,346	15.95	17,657 tok/s
RoPE	237.71 ± 37.67	10,457	17.58	19,413 tok/s
ALiBi	239.67 ± 38.10	10,455	17.71	19,269 tok/s
Relative PE	288.43 ± 24.55	11,517	20.63	16,547 tok/s




Figure 3: Positional encoding perplexity at context 1024. RoPE and ALiBi achieve >40% reduction over baseline. Note that RoPE/ALiBi have larger σ than baseline because their lower mean perplexity makes individual batch variance more visible in absolute units.

5.6 Extrapolation Test

Table 7: Extrapolation test: $L_{\text{train}} = 512$, evaluated at longer lengths without retraining. “FAILS” = CUDA index-out-of-bounds error at runtime: a hard structural failure from a bounded lookup table, not a softmax/NaN collapse.

Model	L=512	L=1024	L=2048	Behaviour
RoPE	266.90 ± 21.45	249.96 ± 28.30	274.15 ± 44.90	Graceful, moderate degradation
ALiBi	278.97 ± 22.10	248.42 ± 19.55	244.26 ± 16.80	Monotonic improvement
Relative PE	320.59 ± 34.12	FAILS	FAILS	Hard structural crash




Figure 4: Extrapolation behaviour of positional encodings. ALiBi uniquely improves with longer context; RoPE degrades gracefully; Relative PE crashes beyond training length due to its bounded embedding table.

5.7 Analysis

Positional encoding is the dominant factor in this study. RoPE and ALiBi reduce perplexity from 406.68 to 237–240, a $>40\%$ reduction, far larger than any attention variant gain ($\approx 5\%$ at best). The mechanism is architectural: learned absolute embeddings require the model to spend parameter budget and training steps discovering that “position 500 is far from position 1.” RoPE and ALiBi encode this prior structurally, freeing the full training budget for semantic learning. At our short 2500-step budget, this head-start is decisive.

ALiBi uniquely improves with longer context. At $L = 2048$ (perplexity 244.26), ALiBi is *better* than at $L = 512$ (278.97). This is the expected behaviour from the ALiBi design: a longer sequence provides more historical token context. Because ALiBi handles positioning via a formula $m_h \cdot |i - j|$ that is defined for any distance, it natively exploits this extra history. The monotonically decreasing σ ($22.10 \rightarrow 19.55 \rightarrow 16.80$) further

confirms that ALiBi becomes *more stable* at longer lengths, not less — the linear penalty smooths out any single token’s outsize influence.

RoPE degrades gracefully but measurably. At $L = 1024$ (249.96), RoPE nearly matches its $L = 512$ performance (266.90). At $L = 2048$ (274.15), it rises by 24 points: a genuine out-of-distribution effect. The rotation angles for positions 513–2048 are mathematically valid but the model was never trained to rely on them, so the attention distributions for those positions are slightly misaligned. This is a fixable limitation (position interpolation or longer training) rather than a fundamental structural failure.

Relative PE fails completely beyond training length due to an index-out-of-bounds error at the embedding table lookup. When $|i - j| \geq 512$, the required table entry does not exist. This is not a NaN or softmax instability, it is a hard runtime crash. No amount of additional training can fix this without expanding the embedding table size and retraining. The design is fundamentally incompatible with length generalisation.

Why RoPE and ALiBi use more memory than baseline. Both methods add $T \times T$ -scale tensors to the forward pass: RoPE precomputes per-position rotation matrices; ALiBi computes a per-head slope bias matrix over the full attention score. These tensors are retained during the backward pass for gradient computation, adding ≈ 2 GB at context 1024.

Relative PE is the slowest and most memory-hungry (11,517 MB, 16,547 tok/s) because it requires an extra einsum computation $(q_i^\top r_{i-j})$ for every query-key pair, on top of the standard attention einsum.

6 Task 4: Convolution + Attention Hybrids

6.1 Motivation

Self-attention captures long-range dependencies but may underutilise the strong local structure in natural language. Adjacent tokens form n-grams and phrase-level patterns that simple local filters can capture extremely efficiently. Depthwise separable Conv1D scans these windows at $\mathcal{O}(T \cdot k \cdot d)$ for kernel size k , complementing the global reach of attention.

6.2 Design 1: Causal Conv1D Before Each Attention Block

A depthwise separable Conv1D is prepended to each attention block, with causality enforced via left-only padding:

$$\hat{x} = x + \text{DepthwiseConv1D}(\text{LN}(x), \text{pad} = (k - 1, 0)) \quad (16)$$

For kernel size 3, token t sees only positions $t - 2, t - 1, t$. The pre-processed token representations then pass through standard attention and FFN sub-blocks.

Bug Detected and Fixed: Symmetric Padding Data Leak

Initial implementation: `nn.Conv1d(embed_dim, embed_dim, 3, padding=1, groups=embed_dim)`

Observed failure: Validation perplexity collapsed to **1.16** — near-perfect, which is physically impossible for a genuine language model.

Root cause: `padding=1` applies one zero on *each* side of the sequence. For kernel size 3, the convolution at token t reads positions $t - 1$, t , and $t + 1$. Position $t + 1$ is a *future token* — the model observes the target it is meant to predict, trivially minimising loss without learning any language structure.

Why it was non-obvious: No runtime error is thrown. Training proceeds normally. The only signal was the suspiciously good perplexity value — which required knowing that WikiText-2 perplexity of 1.16 for a 29M model at 2500 steps is impossible, and that values below ≈ 50 warrant immediate investigation.

Fix: Replace with explicit left-only padding: `xc = F.pad(xc, (2, 0))` before the conv call. Token t now sees exactly $t - 2, t - 1, t$ and no future position.

Result: Perplexity recovered to 345.38 — a realistic value consistent with a model that has genuinely learned to predict text.

General lesson for autoregressive models: Causal correctness must be verified for every architectural component that touches the sequence dimension. Symmetric padding, sliding windows that include future positions, and any form of bi-directional information flow will produce data leakage that manifests as implausibly good training metrics rather than a detectable error.

6.3 Design 2: Interleaved Conv and Attention Blocks

Even-indexed layers (0, 2) apply causal depthwise Conv1D + FFN; odd-indexed layers (1, 3) apply standard MHA + FFN. With 4 total layers, 2 attention and 2 conv layers are used.

Why Interleaved underperforms Design 1. Interleaved Conv achieves 353.12 vs. Design 1's 345.38. Replacing every other attention layer with a local 3-gram convolution halves the number of global attention operations available to route long-distance dependencies. A conv layer at position 0 sees at most 3 tokens and contributes nothing to resolving a coreference between token 1 and token 200. Information from distant context must now travel through only 2 attention layers instead of 4, creating a structural bottleneck that is detectable at WikiText-2 scale.

The design is, however, meaningfully more efficient: 9,818 MB vs. 10,539 MB and 21,846 tok/s vs. 18,797 tok/s, a 9% memory reduction and 16% throughput improvement at a 2% perplexity cost.

6.4 Best-Variant Combined Models

We combine the best attention variant and best positional encoding into both hybrid designs.

Selection rationale:

- **Positional encoding:** Both RoPE (237.71) and ALiBi (239.67) are strong. We test

both. ALiBi is preferred for deployment at variable lengths; RoPE for maximum quality at a fixed training context.

- **Attention variant:** Sliding Window (385.78) outperforms MQA (417.42) on perplexity and is more stable across context lengths (392.31 at 2048 vs. MQA’s 449.39). It is selected as the attention component.

This yields four combined models: Conv+ALiBi+SW, Interleaved+ALiBi+SW, Conv+RoPE+SW, and Interleaved+RoPE+SW.

6.5 Results

Table 8: Convolution hybrid results at context 1024, batch 8. Combined models use Sliding Window attention and the named positional encoding.

Model	Category	PPL ($\mu \pm \sigma$)	Peak GPU (MB)	Epoch (min)	Training Throughput
Baseline	Reference	406.68 ± 17.42	8,346	15.95	17,657 tok/s
Conv+Attention (causal)	Conv Hybrid	345.38 ± 19.80	10,539	18.16	18,797 tok/s
Interleaved Conv	Conv Hybrid	353.12 ± 16.15	9,818	15.62	21,846 tok/s
Conv + ALiBi + SW	Combined	232.49 ± 34.50	8,609	19.97	17,094 tok/s
Interleaved+ALiBi+SW	Combined	233.36 ± 31.25	7,899	17.42	19,597 tok/s
Conv+RoPE+SW	Combined	225.55 ± 29.67	8,390	20.40	16,730 tok/s
Interleaved+RoPE+SW	Combined	234.50 ± 33.10	7,676	17.40	19,614 tok/s




Figure 5: Conv hybrid and combined model perplexity at context 1024. The combined models (rightmost four) achieve >40% improvement over baseline by stacking complementary architectural benefits. Conv+RoPE+SW achieves the overall minimum.

6.6 Analysis

The three components are orthogonally beneficial. Conv+RoPE+SW (225.55) outperforms pure RoPE (237.71), pure Sliding Window (385.78), and pure Conv+Attention (345.38) separately. Each component addresses a distinct gap: Conv1D handles local n-gram patterns that attention handles inefficiently; RoPE provides relative position structure without using training budget; Sliding Window focuses attention capacity on the most contextually relevant window. Because these three modifications target different bottlenecks, their benefits are nearly additive.

Conv+RoPE+SW vs. Conv+ALiBi+SW. At the training context length, RoPE (225.55) beats ALiBi (232.49) by 7 points. For deployment at a fixed context length, RoPE is the better choice. However, based on the extrapolation results in Table 7, ALiBi would be preferred for any deployment scenario where inference context length is variable or exceeds training length, ALiBi’s perplexity continues to improve at $L = 2048$ while RoPE’s rises.

Interleaved design offers the best efficiency tradeoff. Interleaved+RoPE+SW (234.50) uses 7,676 MB at 19,614 tok/s: 9% less memory and 17% higher throughput than Conv+RoPE+SW (8,390 MB, 16,730 tok/s) at a 4% perplexity cost. For memory-constrained deployment, this is a compelling operating point.

7 Bonus: Attention-Free Transformer (AFT) Variants

7.1 Mechanism

The Attention-Free Transformer [10] replaces the dot-product self-attention matrix with element-wise key-gated aggregation. For each position t :

$$Y_t = \sigma(Q_t) \odot \frac{\sum_{s \leq t} \exp(K_s + w_{t,s}) \odot V_s}{\sum_{s \leq t} \exp(K_s + w_{t,s}) + \epsilon} \quad (17)$$

where $w_{t,s}$ is a learned pairwise position bias ($\mathbf{0}$ for AFT-Simple). The query no longer selects individual keys through $q_t^\top k_s$; it gates the already-aggregated value. This eliminates the $T \times T$ attention matrix.

Four variants:

- **AFT-Simple:** $w_{t,s} = 0$. No learned positional bias; $\mathcal{O}(T)$ compute.
- **AFT-Full:** Full $T \times T$ learned bias matrix; reintroduces $\mathcal{O}(T^2)$ storage.
- **AFT-Local:** Learned bias restricted to window $|t - s| \leq w_{\text{size}}$; $\mathcal{O}(Tw)$ bias cost.
- **AFT-Conv:** AFT-Simple augmented with a causal depthwise Conv1D correction, combining linear-time attention with local inductive bias.

7.2 Implementation Note

To prevent numerical overflow in $\exp(K_s + w_{t,s})$, we apply a log-sum-exp stabilisation shift:

Listing 1: Stable AFT core.

```
k_max = torch.max(k, dim=2, keepdim=True)[0]
exp_k = torch.exp(k - k_max)
# For AFT-Full/Local: multiply by exp(relative_bias)
# For AFT-Simple: use exp_k directly
kv_cum = exp_k * v # element-wise
out = torch.sigmoid(q) * (kv_cum / (exp_k + 1e-6))
```

7.3 Experimental Constraints for AFT Evaluation

Why AFT was evaluated at context 256, not 1024

AFT-Simple lacks spatial position constraints. At context 1024 with a 2500-step budget and no positional bias, the model cannot learn reliable position-dependent patterns in time, causing unstable training dynamics. To establish a valid comparison baseline, all four AFT variants were evaluated at context $T = 256$. The Conv+RoPE+SW model was re-run at $T = 256$ to provide a fair same-context reference point. This prevents direct comparison against the $T = 1024$ results in prior sections; AFT would likely score differently at convergence or with longer training.

7.4 Results

Table 9: AFT variant results at context 256, alongside Conv+RoPE+SW re-evaluated at context 256 for a fair comparison. σ across evaluation batches.

Model	PPL ($\mu \pm \sigma$)	Peak GPU (MB)	Training Throughput	Complexity
Conv+RoPE+SW @ ctx 256 (reference)	294.02 $\pm$ 67.41	4,589	16,517 tok/s	$\mathcal{O}(T)$
AFT-Simple	299.89 $\pm$ 47.90	9,500	20,189 tok/s	$\mathcal{O}(T)$
AFT-Full	238.12 $\pm$ 38.45	10,480	15,150 tok/s	$\mathcal{O}(T^2)$
AFT-Local	246.70 $\pm$ 39.22	9,541	18,542 tok/s	$\mathcal{O}(Tw)$
AFT-Conv	209.45 $\pm$ 34.80	9,680	17,812 tok/s	$\mathcal{O}(T)$




Figure 6: AFT variant perplexity at context 256. AFT-Conv achieves the best perplexity in this group by combining linear-time computation with a causal convolution correction pass.

7.5 Analysis

AFT-Simple underperforms the reference model (299.89 vs. 294.02) despite having higher throughput (20,189 vs. 16,517 tok/s). Without any positional information, the model cannot distinguish the distance between tokens, which particularly hurts on WikiText-2 where syntactic agreement and coreference span multiple tokens. The large σ

(47.90) relative to AFT-Full (38.45) confirms this: batch-to-batch perplexity varies heavily because the model has no principled way to weight tokens by their positional relevance.

AFT-Full achieves the best AFT-without-conv perplexity (238.12) but reintroduces $\mathcal{O}(T^2)$ memory (10,480 MB) defeating the primary motivation of attention-free computation. The learned pairwise bias matrix $w_{t,s}$ effectively re-learns a form of attention-like routing, which is why quality recovers, but at the same cost.

AFT-Conv achieves the best overall result in this group (209.45), better than the Conv+RoPE+SW reference at the same context length (294.02). Adding a causal Conv1D correction to the linear AFT-Simple formulation gives the model a local inductive bias that compensates for the missing positional structure in AFT-Simple, while keeping computation $\mathcal{O}(T)$. This confirms the pattern seen in Task 4: local convolutional structure and global context mechanisms are complementary.

8 Task 5: Summary and Comparative Analysis

8.1 Hardware Profiling

The following measurements used a dedicated profiling loop: 5 warmup steps, then 50 timed forward+backward steps with `torch.cuda.reset_peak_memory_stats()` called before the timing window. Batch size 8, context 1024.

Table 10: Hardware profiling results. These numbers reflect isolated per-model measurement; training-log values may differ slightly due to CUDA context fragmentation from prior runs.

Model	Peak GPU Memory	Est. Epoch	Training Throughput
Baseline (MHA)	10,474.9 MB (10.23 GB)	15.95 min	21,400 tok/s
Sliding Window	10,454.0 MB (10.21 GB)	16.51 min	20,677 tok/s
MQA	10,469.5 MB (10.22 GB)	17.73 min	19,251 tok/s
Linear Attn	9,496.5 MB (9.27 GB)	14.96 min	22,813 tok/s
RoPE	10,457.0 MB (10.21 GB)	17.58 min	19,413 tok/s
ALiBi	10,455.0 MB (10.21 GB)	17.71 min	19,269 tok/s
Relative PE	11,517.0 MB (11.25 GB)	20.63 min	16,547 tok/s
Conv+Attention	10,538.5 MB (10.29 GB)	18.16 min	18,797 tok/s
Interleaved Conv	9,818.0 MB (9.59 GB)	15.62 min	21,846 tok/s

8.2 Master Comparison Table

Table 11: Complete results at context 1024. Training throughput is measured *with gradients active* and does not represent inference latency. σ is empirical variance across evaluation batches, not cross-seed variance (single seed throughout).

Model	Category	PPL ($\mu \pm \sigma$)	Peak GPU	Training Throughput	Params
Baseline	Reference	406.68 ± 17.42	8,346 MB	17,657 tok/s	29.2M
Sliding Window	Attention	385.78 ± 13.10	10,566 MB	17,554 tok/s	29.2M
MQA	Attention	417.42 ± 23.05	8,341 MB	17,790 tok/s	28.8M
Linear Attn	Attention	389.27 ± 31.45	7,367 MB	20,602 tok/s	29.2M
RoPE	Pos. Enc.	237.71 ± 37.67	10,457 MB	19,413 tok/s	28.9M
ALiBi	Pos. Enc.	239.67 ± 38.10	10,455 MB	19,269 tok/s	28.9M
Relative PE	Pos. Enc.	288.43 ± 24.55	11,517 MB	16,547 tok/s	29.5M
Conv+Attention	Conv Hybrid	345.38 ± 19.80	10,539 MB	18,797 tok/s	29.2M
Interleaved Conv	Conv Hybrid	353.12 ± 16.15	9,818 MB	21,846 tok/s	29.2M
Conv+ALiBi+SW	Combined	232.49 ± 34.50	8,609 MB	17,094 tok/s	29.0M
Inter+ALiBi+SW	Combined	233.36 ± 31.25	7,899 MB	19,597 tok/s	29.0M
Conv+RoPE+SW	Combined	225.55 ± 29.67	8,390 MB	16,730 tok/s	29.0M
Inter+RoPE+SW	Combined	234.50 ± 33.10	7,676 MB	19,614 tok/s	29.0M




Figure 7: All models ranked by validation perplexity. Positional encoding variants (RoPE, ALiBi) and their combinations dominate the top of the ranking. Error bars show evaluation batch σ .

8.3 Key Findings

1. **Positional encoding dominates quality.** RoPE and ALiBi achieve >40% perplexity reduction vs. baseline, far exceeding attention variants ($\approx 5\%$). The mechanism: relative position encodings provide structural prior knowledge that learned absolute embeddings must discover from data, using up most of a short training budget in the process.

2. **ALiBi is the right choice for variable-length deployment.** It is the only encoding whose perplexity improves monotonically beyond training length (278.97 $\rightarrow$ 244.26 at $L = 2048$), and its evaluation σ *decreases* with context length. More context makes it more stable, not less.
3. **Linear Attention leads on efficiency.** Best training throughput (22,201 tok/s at context 2048) and lowest memory (8,998 MB at context 2048) while baseline needed smaller batch.
4. **The three architectural components are orthogonally beneficial.** Conv+RoPE+SW (225.55) outperforms any single component in isolation, confirming that local conv features, relative positional encoding, and restricted window attention address distinct aspects of language structure.
5. **Interleaved Conv is the best efficiency operating point.** Interleaved+RoPE+SW (234.50): 9% less memory, 17% more throughput than Conv+RoPE+SW at 4% worse perplexity.
6. **AFT-Conv achieves the best quality in the bonus section** (209.45 at context 256), confirming that local convolutional inductive bias is complementary to attention-free linear-time computation, mirroring the finding from Task 4.
7. **Symmetric padding is a non-obvious causal failure mode.** Perplexity 1.16 was produced with no error, only the implausibility of the number itself indicated the bug. Sanity checks on output magnitude are essential in autoregressive model development.

9 Limitations

1. **Short training budget.** 2500 steps processes approximately $2500 \times 8 \times 1024 \approx 20.5\text{M}$ tokens. Full convergence on WikiText-2 typically requires 20,000–50,000 steps at this scale. All perplexity values reflect early-training inductive biases, not converged quality. Relative comparisons are valid (all models trained identically), but absolute numbers are not benchmark-quality.
2. **Single random seed.** Cross-seed variance is not quantified. Due to a fixed computational budget, all models used a single uniform initialisation seed. Differences of a few perplexity points (e.g., MQA 417.42 vs. baseline 406.68) may not be statistically significant across seeds. Only gaps >10 points should be treated as robust findings.
3. **Evaluation σ is batch-variance, not seed-variance.** The σ values throughout this report reflect variation across randomly sampled evaluation batches within a single run — they measure how sensitive each architecture is to the token distribution of a given evaluation batch. They are not equivalent to cross-seed standard deviations.
4. **Throughput is training throughput, not inference latency.** All throughput metrics are measured during the training forward+backward pass with gradients active. True inference latency (autoregressive decoding, no gradient) would be lower and measured differently. The numbers here are appropriate for comparing training

efficiency, not deployment serving cost.

5. **Context 2048 batch size and confounding.** After clearing GPU memory between runs, all models including Baseline were able to run at batch size 4 at context 2048. However, batch 4 introduces noisier gradient estimates than batch 8. The perplexity deterioration at 2048 is therefore partly a batch-noise artefact, not purely architectural. Memory comparisons at 2048 are more trustworthy than perplexity comparisons.
6. **Sliding Window memory overhead is implementation-specific.** Our dense-mask PyTorch implementation allocates the full $T \times T$ score tensor before masking. A sparse kernel or FlashAttention local-window path would have substantially lower memory. The values in Table 5 reflect our implementation, not the architectural minimum.
7. **Small model scale.** At 4 layers, 256 embedding dimension, ≈ 29 M parameters, results may not generalise. Specifically: (a) the dominance of positional encoding over attention variants may diminish at scale with longer training; (b) MQA’s quality penalty shrinks at larger models; (c) the interleaved conv capacity loss may be recoverable with more layers.
8. **No inference benchmarking.** MQA’s primary practical advantage — KV-cache reduction during autoregressive decoding — is not measured here. Training-phase memory and throughput do not capture inference-time benefits.
9. **Context 4096 OOM.** Context length 4096 triggered an OOM allocation failure during the initial attention cache build on the T4 (15 GB VRAM), confirming 2048 as our practical execution ceiling.
10. **Ephemeral checkpoint loss.** Kaggle session timeouts destroyed saved `.pt` checkpoint files in `/kaggle/working/`. All reported metrics were extracted to persistent summary objects before teardown; no data was reconstructed after the fact.

10 Conclusion

This study systematically evaluates long-context Transformer modifications on WikiText-2 under uniform conditions with honest accounting of every experimental constraint. The central finding is unambiguous: **how a model encodes position matters far more than how it attends**, at least at this model scale and training budget. RoPE and ALiBi reduce perplexity by over 40% relative to the baseline, while the best attention variant achieves only $\approx 5\%$. The mechanistic explanation is architectural: relative position encodings provide structural prior knowledge that the model would otherwise infer from data, making far more efficient use of a 2500-step budget.

For practitioners choosing components under compute constraints:

- Use **ALiBi** when deployment context length is variable or exceeds training length.
- Use **RoPE** for maximum quality at a fixed known context length.
- Add **causal Conv1D** before attention blocks for a $\approx 15\%$ perplexity improvement at negligible added memory.

- Use **Sliding Window** when memory is the binding constraint, with the caveat that our dense-mask implementation does not realise theoretical memory savings.
- **Conv+RoPE+SW** is the strongest overall configuration (225.55), fitting within 9 GB.
- **Interleaved+RoPE+SW** is the best efficiency operating point: 9% less memory and 17% more throughput at 4% worse perplexity.

The symmetric padding bug (producing perplexity 1.16 with no runtime error) reinforces a broader principle: in autoregressive models, causal correctness must be *actively verified* at each component. Implausibly good output metrics are a more reliable bug signal than exceptions.

References

- [1] A. Vaswani, N. Shazeer, N. Parmar, J. Uszkoreit, L. Jones, A. N. Gomez, L. Kaiser, and I. Polosukhin. *Attention Is All You Need*. Advances in NeurIPS, 2017.
- [2] N. Shazeer. *Fast Transformer Decoding: One Write-Head is All You Need*. arXiv:1911.02150, 2019.
- [3] A. Katharopoulos, A. Vyas, N. Pappas, and F. Fleuret. *Transformers are RNNs: Fast Autoregressive Transformers with Linear Attention*. ICML, 2020.
- [4] J. Su, Y. Lu, S. Pan, A. Murtadha, B. Wen, and Y. Liu. *RoFormer: Enhanced Transformer with Rotary Position Embedding*. arXiv:2104.09864, 2021.
- [5] O. Press, N. A. Smith, and M. Lewis. *Train Short, Test Long: Attention with Linear Biases Enables Input Length Extrapolation*. ICLR, 2022.
- [6] P. Shaw, J. Uszkoreit, and A. Vaswani. *Self-Attention with Relative Position Representations*. NAACL, 2018.
- [7] I. Beltagy, M. E. Peters, and A. Cohan. *Longformer: The Long-Document Transformer*. arXiv:2004.05150, 2020.
- [8] S. Merity, C. Xiong, J. Bradbury, and R. Socher. *Pointer Sentinel Mixture Models*. ICLR, 2017.
- [9] D. Hendrycks and K. Gimpel. *Gaussian Error Linear Units (GELUs)*. arXiv:1606.08415, 2016.
- [10] S. Zhai, W. Talbott, N. Srivastava, C. Huang, H. Goh, R. Zhang, and J. Susskind. *An Attention Free Transformer*. arXiv:2105.14103, 2021.
- [11] A. Gulati, J. Qin, C. Chiu, N. Parmar, Y. Zhang, J. Yu, W. Han, S. Wang, Z. Zhang, Y. Wu, and R. Pang. *Conformer: Convolution-augmented Transformer for Speech Recognition*. Interspeech, 2020.